

Using target to generate features

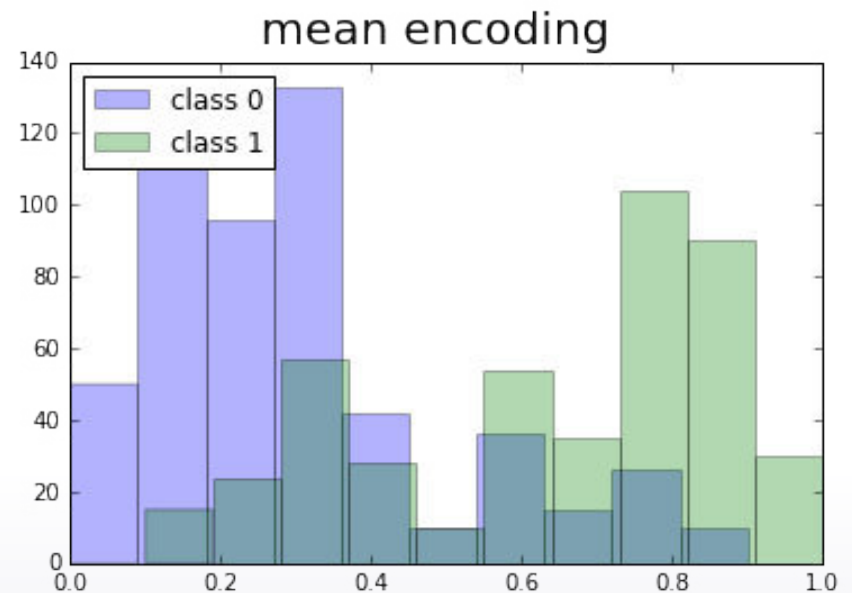
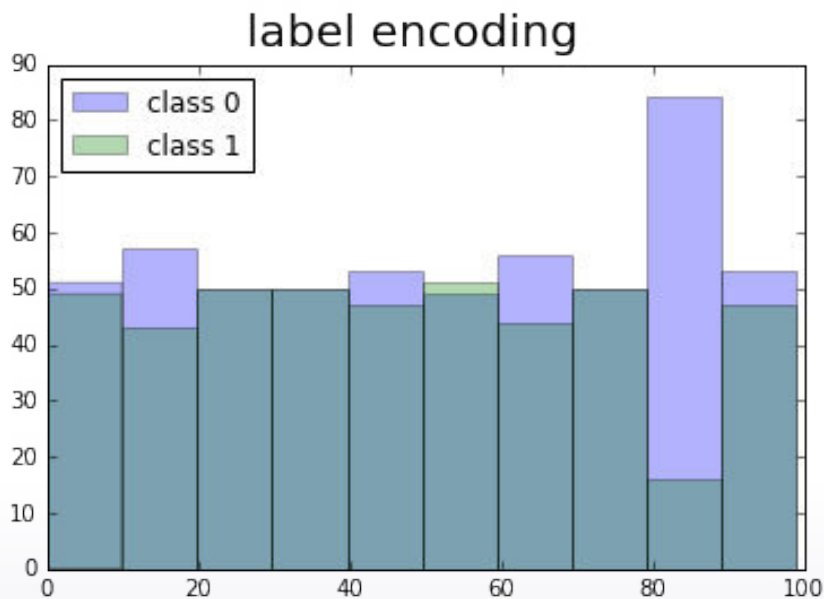
Simple example

- Categorical feature
 - some city
- Binary classification

	feature	feature_label	feature_mean	target
0	Moscow	1	0.4	0
1	Moscow	1	0.4	1
2	Moscow	1	0.4	1
3	Moscow	1	0.4	0
4	Moscow	1	0.4	0
5	Tver	2	0.8	1
6	Tver	2	0.8	1
7	Tver	2	0.8	1
8	Tver	2	0.8	0
9	Klin	0	0.0	0
10	Klin	0	0.0	0
11	Tver	2	0.8	1


Why does it work?

1. Label encoding gives random order. No correlation with target
2. Mean encoding helps to separate zeros from ones



Why does it work?

Reaching a better loss with shorter trees




What will you learn?


- ✓ Construct encodings
- ✓ Correctly validate them
- ✓ Extend them

Indicators of usefulness

Train



Validation



Ways to use target variable

Goods - number of ones in a group,

Bads - number of zeros

- $Likelihood = \frac{Goods}{Goods+Bads} = mean(target)$
- $Weight\ of\ Evidence = \ln\left(\frac{Goods}{Bads}\right) * 100$
- $Count = Goods = sum(target)$
- $Diff = Goods - Bads$

Springleaf example

In [4]:

```
means = X_tr.groupby(col).target.mean()  
train_new[col+'_mean_target'] = train_new[col].map(means)  
val_new[col+'_mean_target'] = val_new[col].map(means)  
  
means
```


Out[4]: VAR_1277

0.0	0.358965
1.0	0.219249
2.0	0.193671
3.0	0.191143
4.0	0.191080
5.0	0.185694


Springleaf example

```
dtrain = xgb.DMatrix(train_new, label=y_tr)
dvalid = xgb.DMatrix(val_new, label=y_val)

evallist = [(dtrain, 'train'), (dvalid, 'eval')]
evals_result3 = {}
model = xgb.train( xgb_par, dtrain, 3000, evals=evallist,
verbose_eval=30, evals_result=evals_result3, early_stopping_rounds=50)
```



Overfit



Train

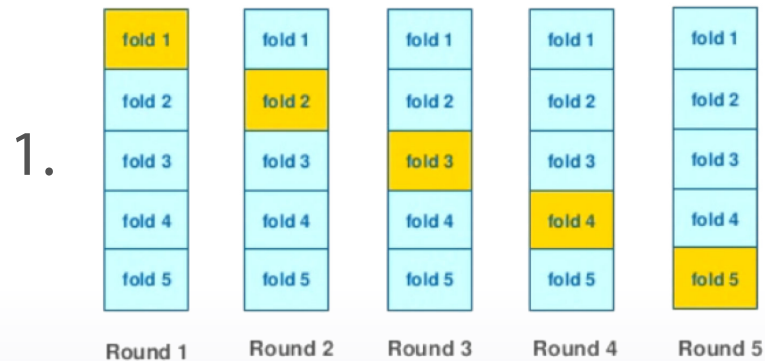
	feature	feature_label	feature_mean	target
8	Tver	2	0	0
9	Klin	0	0	0

Validation

	feature	feature_label	feature_mean	target
10	Klin	0	1	1
11	Tver	2	1	1

Regularization

1. CV loop inside training data;
2. Smoothing;
3. Adding random noise;
4. Sorting and calculating expanding mean.

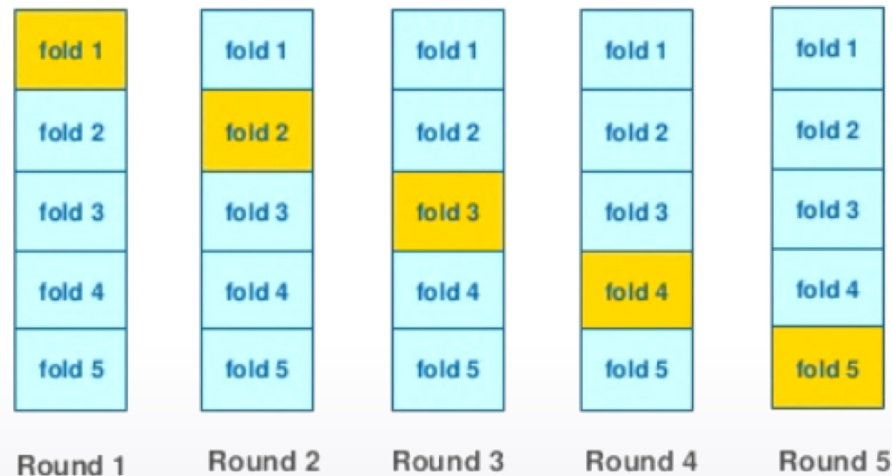


2.
$$\frac{\text{mean}(\text{target}) * \text{nrows} + \text{globalmean} * \alpha}{\text{nrows} + \alpha}$$

Regularization. CV loop

- Robust and intuitive
- Usually decent results with 4-5 folds across different datasets
- Need to be careful with extreme situations like LOO

KFold scheme



Regularization. CV loop

```
y_tr = df_tr['target'].values #target variable
skf = StratifiedKFold(y_tr,5, shuffle=True,random_state=123)

for tr_ind, val_ind in skf:
    X_tr, X_val = df_tr.iloc[tr_ind], df_tr.iloc[val_ind]
    for col in cols: #iterate though the columns we want to encode
        means = X_val[col].map(X_tr.groupby(col).target.mean())
        X_val[col+'_mean_target'] = means
    train_new.iloc[val_ind] = X_val

prior = df_tr['target'].mean() #global mean
train_new.fillna(prior,inplace=True) #fill NaNs with global mean
```

Regularization. CV loop

- Perfect feature for LOO scheme
- Target variable leakage is still present even for KFold scheme

Leave-one-out

	feature	feature_mean	target
0	Moscow	0.50	0
1	Moscow	0.25	1
2	Moscow	0.25	1
3	Moscow	0.50	0
4	Moscow	0.50	0

Regularization.Smoothing

- Alpha controls the amount of regularization
- Only works together with some other regularization method

$$\frac{\text{mean}(\text{target}) * \text{nrows} + \text{globalmean} * \text{alpha}}{\text{nrows} + \text{alpha}}$$

Regularization. Noise

- Noise degrades the quality of encoding
- How much noise should we add?
- Usually used together with LOO

Regularization. Expanding mean

- Least amount of leakage
- No hyper parameters
- Irregular encoding quality
- Built - in in CatBoost

```
cumsum = df_tr.groupby(col)['target'].cumsum() - df_tr['target']  
cumcnt = df_tr.groupby(col).cumcount()  
train_new[col+'_mean_target'] = cumsum/cumcnt
```

Regularization. Conclusion

- There are a lot ways to regularize mean encodings
- Unending battle with target variable leakage
- CV loop or Expanding mean for practical tasks

Generalizations and extensions

- Using target variable in different tasks. Regression, multiclass
- Domains with many-to-many relations
- Timeseries
- Encoding interactions and numerical features

Regression and multiclass

- More statistics for regression tasks. Percentiles, std, distribution bins.
- Introducing new information for one vs all classifiers in multi class tasks

Many-to-many relations

- Cross product of entities
- Statistics from vectors

User_id	APPS	Target
10	APP1; APP2; APP3	0
11	APP4; APP1	1
12	APP2	1
100	APP3; APP9	0

LONG REPRESENTATION

User_id	APP_id	Target
10	APP1	0
10	APP2	0
10	APP3	0
11	APP4	1
11	APP1	1


Time series

- Time structure allows us to make a lot of complicated features.
- Rolling statistics of target variable

Day	User	Spend	Amount	Prev_user	Prev_spend_avg
1	101	FOOD	2.0	0.0	0.0
1	101	GAS	4.0	0.0	0.0
1	102	FOOD	3.0	0.0	0.0
2	101	GAS	4.0	6.0	4.0
2	101	TV	8.0	6.0	0.0
2	102	FOOD	2.0	3.0	2.5

Interactions and numerical features

- Analyzing fitted model
- Binning numeric and selecting interactions



Amazon.com - Employee Access Challenge Competition

Your most recent submission

Name	Submitted	Wait time	Execution time	Score
cat_boost1.csv	a few seconds ago	0 seconds	0 seconds	0.91581

Complete

[Jump to your position on the leaderboard ▼](#)

Your most recent submission


Name	Submitted	Wait time	Execution time	Score
lgb1.csv	just now	0 seconds	0 seconds	0.87209

Complete

[Jump to your position on the leaderboard ▼](#)

Correct validation reminder

- Local experiments:
 - Estimate encodings on X_{tr}
 - Map them to X_{tr} and X_{val}
 - Regularize on X_{tr}
 - Validate model on X_{tr}/X_{val} split
- Submission:
 - Estimate encodings on whole Train data
 - Map them to Train and Test
 - Regularize on Train
 - Fit on Train

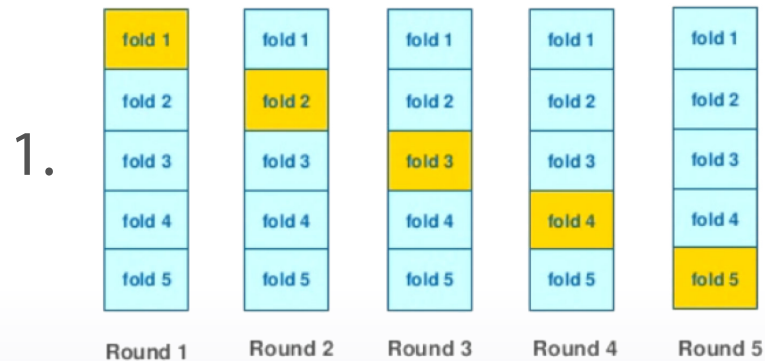


End

- Main advantages:
 - Compact transformation of categorical variables
 - Powerful basis for feature engineering
- Disadvantages:
 - Need careful validation, there a lot of ways to overfit
 - Significant improvements only on specific datasets

Regularization

1. CV loop inside training data;
2. Smoothing;
3. Adding random noise;
4. Sorting and calculating expanding mean.

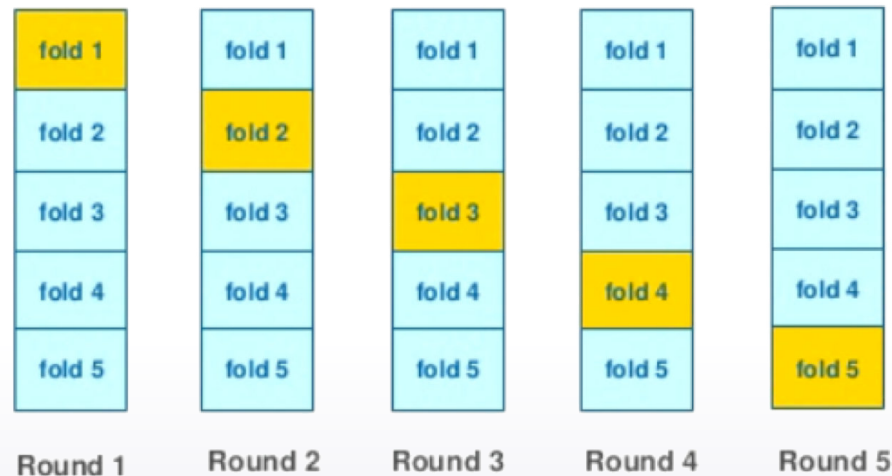


2.
$$\frac{\text{mean}(\text{target}) * \text{nrows} + \text{globalmean} * \alpha}{\text{nrows} + \alpha}$$

Regularization. CV loop

- Robust and intuitive
- Usually decent results with 4-5 folds across different datasets
- Need to be careful with extreme situations like LOO

KFold scheme



Regularization. CV loop

```
y_tr = df_tr['target'].values #target variable
skf = StratifiedKFold(y_tr,5, shuffle=True,random_state=123)

for tr_ind, val_ind in skf:
    X_tr, X_val = df_tr.iloc[tr_ind], df_tr.iloc[val_ind]
    for col in cols: #iterate though the columns we want to encode
        means = X_val[col].map(X_tr.groupby(col).target.mean())
        X_val[col+'_mean_target'] = means
    train_new.iloc[val_ind] = X_val

prior = df_tr['target'].mean() #global mean
train_new.fillna(prior,inplace=True) #fill NaNs with global mean
```

Regularization. CV loop

- Perfect feature for LOO scheme
- Target variable leakage is still present even for KFold scheme

Leave-one-out

	feature	feature_mean	target
0	Moscow	0.50	0
1	Moscow	0.25	1
2	Moscow	0.25	1
3	Moscow	0.50	0
4	Moscow	0.50	0

Regularization.Smoothing

- Alpha controls the amount of regularization
- Only works together with some other regularization method

$$\frac{\text{mean}(\text{target}) * \text{nrows} + \text{globalmean} * \text{alpha}}{\text{nrows} + \text{alpha}}$$

Regularization. Noise

- Noise degrades the quality of encoding
- How much noise should we add?
- Usually used together with LOO

Regularization. Expanding mean

- Least amount of leakage
- No hyper parameters
- Irregular encoding quality
- Built - in in CatBoost

```
cumsum = df_tr.groupby(col)['target'].cumsum() - df_tr['target']  
cumcnt = df_tr.groupby(col).cumcount()  
train_new[col+'_mean_target'] = cumsum/cumcnt
```

Regularization. Conclusion

- There are a lot ways to regularize mean encodings
- Unending battle with target variable leakage
- CV loop or Expanding mean for practical tasks